

NAME:- SHAURYA PRAKASH SHAH

ROLL:- 001811001025

ASSIGNMENT – 5

IT 4TH YEAR 1ST SEM

IPYNB Notebook Link:-

<https://colab.research.google.com/drive/13jQE77EuFZxJsO72X59Od4jMGGmmBskS?usp=sharing>

GITHUB Link:- <https://github.com/shauryashah/ML-Lab-Assignments.git>

1. MOUNTAIN CAR - REINFORCEMENT LEARNING

CODE

```
import gym
import numpy as np
import matplotlib.pyplot as plt
import networkx as nx
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Reshape, Conv2D, Dense, Flatten, BatchNormalization, Dropout, MaxPooling2D
from tensorflow.keras.optimizers import Adam, SGD

from rl.agents import DQNAgent
from rl.policy import EpsGreedyQPolicy
from rl.memory import SequentialMemory

def plot_average_reward(reward_list, ave_reward_list):
    plt.plot(np.arange(len(reward_list)), reward_list, label='Episode Reward')
    plt.plot(100*(np.arange(len(ave_reward_list)) + 1), ave_reward_list, label='Average Reward')
    plt.legend(loc='upper left')
    plt.xlabel('Episodes')
    plt.title('Average Reward vs Episodes')
    plt.show()
```

```

env = gym.make('MountainCar-v0')
env.reset()

print('State space: ', env.observation_space)
print('Action space: ', env.action_space)
print(env.observation_space.low)
print(env.observation_space.high)

%%time

learning = 0.1
discount = 0.95
epsilon = 0.5
min_eps = 0.0
episodes = 5000
epsilon_decay = (epsilon-min_eps)/episodes
reward_list = []
ave_reward_list = []
win_count = 0
cumu_timesteps = 0

discrete_obs_size = [20]*len(env.observation_space.high)
discrete_obs_window = (env.observation_space.high-
env.observation_space.low)/discrete_obs_size
q_table = np.random.uniform(low=-
2, high=0, size=(discrete_obs_size+[env.action_space.n]))

def get_discrete_state(state):
    discrete_state = (state-
env.observation_space.low)/discrete_obs_window
    return tuple(discrete_state.astype(np.int))

for episode in range(episodes):
    discrete_state = get_discrete_state(env.reset())
    tot_reward = 0
    done=False
    while not done:

        cumu_timesteps+=1
        if np.random.random() > epsilon:
            action = np.argmax(q_table[discrete_state])
        else:
            action = np.random.randint(0, env.action_space.n)

        new_state, reward, done, _ = env.step(action)
        new_discrete_state = get_discrete_state(new_state)

```

```

    if not done:
        max_future_q = np.max(q_table[new_discrete_state])
        current_q = q_table[discrete_state + (action,)]
        new_q = (1-
learning)*current_q + learning*(reward + discount*max_future_q)
        q_table[discrete_state + (action,)] = new_q

    elif new_state[0] >= env.goal_position:
        win_count+=1
        q_table[discrete_state + (action,)] = 0

    tot_reward+=reward
    discrete_state = new_discrete_state

    if epsilon > min_eps:
        epsilon-=epsilon_decay
    reward_list.append(tot_reward)

    if (episode+1) % 100 == 0:
        ave_reward = np.mean(reward_list[episode-99:])
        ave_reward_list.append(ave_reward)


    if (episode+1) % 500 == 0:
        print('Episode {} Average Reward: {}'.format(episode+1, ave_rewar
d))
env.close()

```

```

Episode 500 Average Reward: -200.0
Episode 1000 Average Reward: -200.0
Episode 1500 Average Reward: -199.2
Episode 2000 Average Reward: -197.14
Episode 2500 Average Reward: -195.4
Episode 3000 Average Reward: -186.4
Episode 3500 Average Reward: -168.82
Episode 4000 Average Reward: -159.06
Episode 4500 Average Reward: -177.29
Episode 5000 Average Reward: -161.02
CPU times: user 1min 52s, sys: 4.76 s, total: 1min 57s
Wall time: 1min 51s

```



2. MOUNTAIN CAR – DEEP REINFORCEMENT LEARNING

CODE

```
#DEFINING ACTION AND STATE SIZE
nb_actions = 3
nb_states = 2

#DEFINE DEEP LEARNING MODEL
model = Sequential()
model.add(Flatten(input_shape=(1,) + env.observation_space.shape))
model.add(Dense(128, activation='relu'))
model.add(Dense(512, activation="relu"))
model.add(Dropout(0.5))
model.add(Dense(3, activation="relu"))
model.summary()

#INITIALISE AGENT WITH EPSILON GREDDY POLICY
policy = EpsGreedyQPolicy()
memory = SequentialMemory(limit=5000, window_length=1)
agent = DQNAgent(model=model, memory=memory, policy=policy, nb_actions=
nb_actions,
                    nb_steps_warmup=500, target_model_update=1e-2)
agent.compile(Adam(lr=1e-3), metrics=['mse'])

#AGENT IS TRAINED ON ENIRONMENT
agent.fit(env, nb_steps=50000, visualize=False, verbose=1, nb_max_episo
de_steps=1000)
```

```
#AGENT IS TESTED FOR 10 EPISODES
agent.test(env, nb_episodes=10, nb_max_episode_steps=1000, visualize=False)
```

OUTPUT

Model: "sequential_14"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 2)	0
dense_33 (Dense)	(None, 128)	384
dense_34 (Dense)	(None, 512)	66048
dropout_10 (Dropout)	(None, 512)	0
dense_35 (Dense)	(None, 3)	1539

Total params: 67,971
 Trainable params: 67,971
 Non-trainable params: 0

```

Training for 50000 steps ...
Interval 1 (0 steps performed)
/usr/local/lib/python3.7/dist-packages/keras/engine/training_v1.py:2079: UserWarning
  updates=self.state_updates,
10000/10000 [=====] - 118s 12ms/step - reward: -1.0000
50 episodes - episode_reward: -200.000 [-200.000, -200.000] - loss: 0.500 - mse: 0.3

Interval 2 (10000 steps performed)
10000/10000 [=====] - 122s 12ms/step - reward: -1.0000
50 episodes - episode_reward: -200.000 [-200.000, -200.000] - loss: 0.500 - mse: 0.3

Interval 3 (20000 steps performed)
10000/10000 [=====] - 123s 12ms/step - reward: -1.0000
50 episodes - episode_reward: -200.000 [-200.000, -200.000] - loss: 0.500 - mse: 0.3

Interval 4 (30000 steps performed)
10000/10000 [=====] - 120s 12ms/step - reward: -1.0000
50 episodes - episode_reward: -200.000 [-200.000, -200.000] - loss: 0.500 - mse: 0.3

Interval 5 (40000 steps performed)
10000/10000 [=====] - 124s 12ms/step - reward: -1.0000
done, took 606.494 seconds
<keras.callbacks.History at 0x7f5a9be0be50>

```

```
Testing for 10 episodes ...
Episode 1: reward: -200.000, steps: 200
Episode 2: reward: -200.000, steps: 200
Episode 3: reward: -200.000, steps: 200
Episode 4: reward: -200.000, steps: 200
Episode 5: reward: -200.000, steps: 200
Episode 6: reward: -200.000, steps: 200
Episode 7: reward: -200.000, steps: 200
Episode 8: reward: -200.000, steps: 200
Episode 9: reward: -200.000, steps: 200
Episode 10: reward: -200.000, steps: 200
<keras.callbacks.History at 0x7f5a9bafca90>
```

3. ROULETTE - REINFORCEMENT LEARNING

CODE

```
env = gym.make('Roulette-v0')
env.reset()
print('State space: ', env.observation_space)
print('Action space: ', env.action_space)

learning = 0.1
discount = 0.95
epsilon = 0.5
min_eps = 0.0
episodes = 50000
epsilon_decay = (epsilon-min_eps)/episodes
reward_list = []
ave_reward_list = []

q_table = np.random.randn(env.observation_space.n, env.action_space.n)
for episode in range(episodes):
    state = env.reset()
    tot_reward = 0
    done=False
    while not done:
        if np.random.random() > epsilon:
            action = np.argmax(q_table[state,:])
        else:
            action = np.random.randint(0, env.action_space.n)

        new_state, reward, done, _ = env.step(action)
        max_future_q = np.max(q_table[new_state, :])
        current_q = q_table[state, action]
        new_q = (1-
learning)*current_q + learning*(reward + discount*max_future_q)
        q_table[state, action] = new_q
```

```

        tot_reward+=reward
        state = new_state

    if epsilon > min_eps:
        epsilon-=epsilon_decay
        reward_list.append(tot_reward)

    if (episode+1) % 100 == 0:
        ave_reward = np.mean(reward_list[episode-99:])
        ave_reward_list.append(ave_reward)
        #reward_list = []

    if (episode+1) % 500 == 0:
        print('Episode {} Average Reward: {}'.format(episode+1, ave_reward))
d))
env.close()


```

OUTPUT

```

Episode 35500 Average Reward: -1.22
Episode 36000 Average Reward: -0.37
Episode 36500 Average Reward: -0.28
Episode 37000 Average Reward: -0.76
Episode 37500 Average Reward: -0.23
Episode 38000 Average Reward: -0.32
Episode 38500 Average Reward: -0.53
Episode 39000 Average Reward: -0.11
Episode 39500 Average Reward: -0.11
Episode 40000 Average Reward: -0.17
Episode 40500 Average Reward: -0.07
Episode 41000 Average Reward: -0.03
Episode 41500 Average Reward: -0.12
Episode 42000 Average Reward: 0.0
Episode 42500 Average Reward: -0.05
Episode 43000 Average Reward: -0.04
Episode 43500 Average Reward: 0.0
Episode 44000 Average Reward: 0.05
Episode 44500 Average Reward: -0.05
Episode 45000 Average Reward: -0.26
Episode 45500 Average Reward: -0.06
Episode 46000 Average Reward: -0.01
Episode 46500 Average Reward: -0.01
Episode 47000 Average Reward: -0.02
Episode 47500 Average Reward: -0.01
Episode 48000 Average Reward: 0.01
Episode 48500 Average Reward: -0.01
Episode 49000 Average Reward: 0.0
Episode 49500 Average Reward: 0.0
Episode 50000 Average Reward: 0.0
CPU times: user 42.8 s, sys: 4.92 s, total: 47.7 s
Wall time: 43 s

```



4. ROULETTE – DEEP REINFORCEMENT LEARNING

CODE

```
env = gym.make('Roulette-v0')
print(env.observation_space)
print(env.observation_space.shape)
print(env.action_space)
print(env.action_space.shape)

#DEFINE ACTION AND STATES
nb_actions = 38
nb_states = 1

#DEFINE DEEP LEARNING MODEL
model = Sequential()
model.add(Dense(128, input_dim=nb_states))
model.add(Dense(256, activation="relu"))
model.add(Dropout(0.5))
model.add(Dense(38, activation="relu"))
model.summary()

#DEFINE AGENT WITH EPSILON GREEDY POLICY
policy = EpsGreedyQPolicy()
memory = SequentialMemory(limit=5000, window_length=1)
agent = DQNAgent(model=model, memory=memory, policy=policy, nb_actions=
nb_actions,
                  nb_steps_warmup=500, target_model_update=1e-2)
agent.compile(Adam(lr=1e-3), metrics=['mse'])
```



```
#TRAIN AGENT ON ENVIRONMENT
agent.fit(env, nb_steps=50000, visualize=False, verbose=1, nb_max_episode_steps=1000)
```

```
#TEST AGENT ON 10 EPISODES
agent.test(env, nb_episodes=10, nb_max_episode_steps=1000, visualize=False)
```

OUTPUT

Model: "sequential_5"

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 128)	256
dense_9 (Dense)	(None, 256)	33024
dropout_2 (Dropout)	(None, 256)	0
dense_10 (Dense)	(None, 38)	9766

```
=====
Total params: 43,046
Trainable params: 43,046
Non-trainable params: 0
=====
```

```
Training for 50000 steps ...
Interval 1 (0 steps performed)
  53/10000 [.....] - ETA: 9s - reward: -0.1509 /usr/local
updates=self.state_updates,
10000/10000 [=====] - 97s 10ms/step - reward: 0.0032
109 episodes - episode_reward: 0.394 [-98.000, 245.000] - loss: 17.272 - mse: 0.909

Interval 2 (10000 steps performed)
10000/10000 [=====] - 104s 10ms/step - reward: 0.0165
119 episodes - episode_reward: 1.487 [-96.000, 132.000] - loss: 16.026 - mse: 0.843

Interval 3 (20000 steps performed)
10000/10000 [=====] - 103s 10ms/step - reward: 0.0272
114 episodes - episode_reward: 2.254 [-96.000, 136.000] - loss: 16.692 - mse: 0.879

Interval 4 (30000 steps performed)
10000/10000 [=====] - 101s 10ms/step - reward: -0.0149
114 episodes - episode_reward: -1.386 [-92.000, 245.000] - loss: 16.559 - mse: 0.872

Interval 5 (40000 steps performed)
10000/10000 [=====] - 101s 10ms/step - reward: 0.0540
done, took 506.263 seconds
<keras.callbacks.History at 0x7f5a9cbaed90>
```

```
Testing for 10 episodes ...
Episode 1: reward: -26.000, steps: 100
Episode 2: reward: 11.000, steps: 100
Episode 3: reward: -63.000, steps: 100
Episode 4: reward: -26.000, steps: 100
Episode 5: reward: 11.000, steps: 100
Episode 6: reward: 11.000, steps: 100
Episode 7: reward: -26.000, steps: 100
Episode 8: reward: -63.000, steps: 100
Episode 9: reward: -26.000, steps: 100
Episode 10: reward: 85.000, steps: 100
<keras.callbacks.History at 0x7f5a9b877f10>
```

5. CAR RACING – REINFORCEMENT LEARNING

Car Racing was not implemented using q learning as the q table would be very large of the order of action space length * $256^{(96*96*3)}$. Since action space is not discrete but continuous, hence the dimensions would be greater(equal to no of buckets required). **Hence, the RAM on Colab kept crashing while trying to train.** Also each state of the car racing game is a snapshot of the current status of the game. Normal q learning methods are not good enough to train on this. A CNN would have much better results.

6. CAR RACING – DEEP REINFORCEMENT LEARNING

CODE

```
env = gym.make('CarRacing-v0')
print(env.observation_space)
print(env.observation_space.shape)
print(env.action_space)
print(env.action_space.shape)
```

```
#DEFINE WRAPPER CLASS FOR CARRACING-V0 TO MAKE ACTIONS DISCRETE
```

```
class CarRacingDiscrit:
```

```
    def __init__(self):
```

```
        self.env = gym.make('CarRacing-v0')
```

```
        self.action_space = 10*10*10
```

```
        self.observation_space = 96*96*3
```

```
    def step(self, action):
```

```
        v1 = int(      action      ) % 10
```

```
        v2 = int( int(action) / 10 ) % 10
```

```
        v3 = int( int(action) / 100 ) % 10
```

```
        v1 = ( v1 - 5 ) / 5
```

```
        v2 = ( v2      ) / 10
```

```
        v3 = ( v3      ) / 10
```

```
        state, reward, done, info = self.env.step([v1, v2, v3])
```

```
        return state, reward, done, info
```

```
    def seed(self, s):
```

```
        return env.seed(s)
```

```
    def reset(self):
```

```
        return self.env.reset()
```

```
    def render(self):
```

```
        return self.env.render()
```

```
    def close(self):
```

```
        return self.env.close()
```

```

env = CarRacingDiscrit()

nb_actions = 10*10*10

print(env.observation_space)

print(env.action_space)


#DEFINE DEEP LEARNING MODEL

model = Sequential()

model.add(Reshape((96, 96, 3), input_shape=(1, 96, 96, 3)))

model.add(BatchNormalization())

model.add(Conv2D(filters=32, kernel_size=(3, 3), activation="relu"))

model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(BatchNormalization())

model.add(Conv2D(filters=64, kernel_size=(3, 3), activation="relu"))

model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())

model.add(Dense(192, activation="relu"))

model.add(Dropout(0.5))

model.add(Dense(1000, activation="relu"))

model.summary()


#DEFINE AGENT WITH EPSILON GREEDY POLICY

policy = EpsGreedyQPolicy()

memory = SequentialMemory(limit=5000, window_length=1)

agent = DQNAgent(model=model, memory=memory, policy=policy, nb_actions=
nb_actions,

                    nb_steps_warmup=500, target_model_update=1e-2)

agent.compile(Adam(lr=1e-3), metrics=['mse'])


#TRAIN AGENT ON ENVIRONMENT

agent.fit(env, nb_steps=10000, visualize=False, verbose=1, nb_max_episo
de_steps=1000)

```

```
#TEST AGENT FOR 10 EPISODES
```

```
agent.test(env, nb_episodes=10, nb_max_episode_steps=1000, visualize=False)
```

OUTPUT

Model: "sequential"

Layer (type)	Output Shape	Param #
reshape (Reshape)	(None, 96, 96, 3)	0
batch_normalization (Batch Normalization)	(None, 96, 96, 3)	12
conv2d (Conv2D)	(None, 94, 94, 32)	896
max_pooling2d (MaxPooling2D)	(None, 47, 47, 32)	0
batch_normalization_1 (Batch Normalization)	(None, 47, 47, 32)	128
conv2d_1 (Conv2D)	(None, 45, 45, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 22, 22, 64)	0
flatten (Flatten)	(None, 30976)	0
dense (Dense)	(None, 192)	5947584
dropout (Dropout)	(None, 192)	0
dense_1 (Dense)	(None, 1000)	193000

```
=====  
Total params: 6,160,116  
Trainable params: 6,160,046  
Non-trainable params: 70
```

```
Training for 10000 steps ...
Track generation: 1163..1458 -> 295-tiles track
Interval 1 (0 steps performed)
/usr/local/lib/python3.7/dist-packages/keras/engine/training_v1.py:2079: UserWarning:
  updates=self.state_updates,
1000/10000 [==>.....] - ETA: 10:01 - reward: -0.0796Track ge
2000/10000 [=====>.....] - ETA: 12:12 - reward: -0.0319Track ge
3000/10000 [=====>.....] - ETA: 11:39 - reward: -0.0324Track ge
4000/10000 [=====>.....] - ETA: 10:27 - reward: -0.0432Track ge
5000/10000 [=====>.....] - ETA: 8:58 - reward: -0.0445Track gen
6000/10000 [=====>.....] - ETA: 7:18 - reward: -0.0433Track gen
7000/10000 [=====>.....] - ETA: 5:32 - reward: -0.0421Track gen
8000/10000 [=====>.....] - ETA: 3:43 - reward: -0.0421Track gen
9000/10000 [=====>.....] - ETA: 1:52 - reward: -0.0447Track gen
10000/10000 [=====>.....] - 1157s 113ms/step - reward: -0.0419
done, took 1159.031 seconds
<keras.callbacks.History at 0x7f8f803a4490>
```

```
Testing for 10 episodes ...
Track generation: 1186..1490 -> 304-tiles track
Episode 1: reward: -83.498, steps: 1000
Track generation: 1177..1476 -> 299-tiles track
Episode 2: reward: -83.221, steps: 1000
Track generation: 1108..1389 -> 281-tiles track
Episode 3: reward: -82.143, steps: 1000
Track generation: 1004..1265 -> 261-tiles track
Episode 4: reward: -80.769, steps: 1000
Track generation: 1111..1393 -> 282-tiles track
Episode 5: reward: -82.206, steps: 1000
Track generation: 1120..1413 -> 293-tiles track
Episode 6: reward: -82.877, steps: 1000
Track generation: 1132..1419 -> 287-tiles track
Episode 7: reward: -82.517, steps: 1000
Track generation: 1184..1493 -> 309-tiles track
Episode 8: reward: -83.766, steps: 1000
Track generation: 1041..1310 -> 269-tiles track
Episode 9: reward: -81.343, steps: 1000
Track generation: 1053..1326 -> 273-tiles track
Episode 10: reward: -81.618, steps: 1000
<keras.callbacks.History at 0x7f8e93fb3ad0>
```

2. USER INPUT GRAPH

CODE

ASSUMING A GRAPH IN THE SHAPE OF A GRID WITH MOVEMENT ALLOWED IN ALL DIRECTIONS EXCEPT ALONG THE DIAGONALS

```
import random, math, time

import numpy as np

from keras.models import Sequential

from keras.layers import *

from tensorflow.keras.optimizers import *


import matplotlib

#matplotlib.use("Agg")

import matplotlib.pyplot as plt

from matplotlib.image import imread

from matplotlib import rc, animation

from IPython import display

from IPython.display import HTML

%matplotlib inline


#DEFINING AN ENVIRONMENT FOR A USER INPUT GRAPH

class Environment:

    def __init__(self, grid_size):

        self.grid_size = grid_size


        self.cat = imread('start.png')

        self.mouse = imread('dest.jpg')
```

```

        #self.confetti = imread('https://image.ibb.co/ganuAA/tom-and-jerry.png')

        self.dim = 1.5

        self.rewards = []

    def _update_state(self, action):
        state = self.state

        # 0 = left
        # 1 = right
        # 2 = down
        # 3 = up

        fy, fx, py, px = state
        old_d = abs(fx - px) + abs(fy - py)

        if action == 0:
            if px > 0:
                px -= 1
        if action == 1:
            if px < self.grid_size-1:
                px += 1
        if action == 2:
            if py > 0:
                py -= 1
        if action == 3:
            if py < self.grid_size-1:
                py += 1

        new_d = abs(fx - px) + abs(fy - py)
        self.d = old_d-new_d
        self.time = self.time - 1

```



```

        return np.array([fy, fx, py, px])

def _get_reward(self):
    fruit_y, fruit_x, player_y, player_x = self.state
    if fruit_x == player_x and fruit_y == player_y: return 1
    if self.d == 1: return 1
    if self.d == 0: return -1
    if self.d == -1: return -1

def _is_over(self):
    fruit_y, fruit_x, player_y, player_x = self.state
    if self.time == 0: return True
    if fruit_x == player_x and fruit_y == player_y: return True
    return False

def step(self, action):
    self.state = self._update_state(action)
    reward = self._get_reward()
    self.rewards.append(reward)
    game_over = self._is_over()
    return self.state, reward, game_over

def render(self):
    # Note: there's no promises of efficiency with this method
    # If things are slow, remove it

    im_size = (self.grid_size,)*2
    state = self.state

    self.fig = plt.figure(figsize=(8, 6), dpi=80)
    self.ax = self.fig.add_subplot(111)

```

```

self.ax.clear()
self.ax.set_ylim((-1, self.grid_size))
self.ax.set_xlim((-1, self.grid_size))
#self.ax.axis('off') # uncomment to turn off axes
self.ax.get_xaxis().set_ticks(range(self.grid_size))
self.ax.get_yaxis().set_ticks(range(self.grid_size))

xc = state[2]
yc = state[3]
xm = state[0]
ym = state[1]

if state[0] == state[2] and state[1] == state[3]:
    self.ax.imshow(self.cat,
                    extent=(-1, self.grid_size,
                             -1, self.grid_size))
else:
    self.ax.imshow(self.mouse,
                    extent=(xm-self.dim/4, xm+self.dim/4,
                             ym-self.dim/4, ym+self.dim/4))
    self.ax.imshow(self.cat,
                    extent=(xc-self.dim/4, xc+self.dim/4,
                             yc-self.dim/4, yc+self.dim/4))
self.fig.canvas.draw()
return np.array(self.fig.canvas.renderer._renderer)

def reset(self, deterministic=True):
    if deterministic:
        # this is an easier environment setup
        fruit_x = 0

```

```

        fruit_y = 0
        player_x = self.grid_size - 1
        player_y = self.grid_size - 1
        time = self.grid_size*2
    else:
        generated = False
        while not generated\
            or abs(fruit_x - player_x) + abs(fruit_y - player_y) < self.grid_
size/2:
            fruit_x = np.random.randint(0, self.grid_size-1)
            fruit_y = np.random.randint(0, self.grid_size-1)
            player_x = np.random.randint(0, self.grid_size-1)
            player_y = np.random.randint(0, self.grid_size-1)
            time = abs(fruit_x - player_x) + abs(fruit_y - player_y)
            time *= 2
            generated = True

    self.time = time
    self.d = 0
    self.state = np.asarray([fruit_y, fruit_x, player_y, player_x])

    return self.state

- """
- This runs the environment using random actions
- """

print('Setting up environment')
env = Environment(5)
num_episodes = 1 # number of games we want the agent to play
env.reset()

```

```

frames = []

RENDER = True

print('Running random simulation')

for episode in range(num_episodes):
    print('Resetting environment')
    s = env.reset() # Initial state
    while True:
        a = np.random.choice(range(4)) # choose a random action
        s_, r, done = env.step(a) # apply random action

        if RENDER:
            fig = env.render()
            plt.imshow(fig)
            plt.show()
            frames.append(fig)

        if done:
            break

```

REINFORCEMENT LEARNING

```

%%time

learning = 0.1
discount = 0.95
epsilon = 0.5
min_eps = 0.0
episodes = 5000
epsilon_decay = (epsilon-min_eps)/episodes
reward_list = []
ave_reward_list = []

```

```

win_count = 0
cumu_timesteps = 0

discrete_obs_size = [5]*4
q_table = np.random.uniform(low=-
2, high=0, size=(discrete_obs_size+[4]))
print(q_table.shape)

def get_discrete_state(state):
    discrete_state = (state-
env.observation_space.low)/discrete_obs_window
    return tuple(discrete_state.astype(np.int))

for episode in range(episodes):
    state = tuple(env.reset().astype(np.int))
    tot_reward = 0
    done=False
    while not done:

        cumu_timesteps+=1
        if np.random.random() > epsilon:
            action = np.argmax(q_table[state])
        else:
            action = np.random.randint(0, 4)

        new_state, reward, done = env.step(action)
        new_state = tuple(new_state.astype(np.int))
        if not done:
            max_future_q = np.max(q_table[new_state])
            current_q = q_table[state + (action,)]
            new_q = (1-
learning)*current_q + learning*(reward + discount*max_future_q)

```

```

        q_table[state + (action,)] = new_q

    elif done:

        win_count+=1

    tot_reward+=reward

    state = new_state

    if epsilon > min_eps:

        epsilon-=epsilon_decay

    reward_list.append(tot_reward)

    if (episode+1) % 100 == 0:

        ave_reward = np.mean(reward_list[episode-99:])

        ave_reward_list.append(ave_reward)

    if (episode+1) % 500 == 0:

        print('Episode {} Average Reward: {}'.format(episode+1, ave_reward))

print(win_count)


```

OUTPUT


```

(5, 5, 5, 5, 4)
Episode 500 Average Reward: 4.56
Episode 1000 Average Reward: 5.3
Episode 1500 Average Reward: 5.65
Episode 2000 Average Reward: 5.48
Episode 2500 Average Reward: 5.63
Episode 3000 Average Reward: 5.79
Episode 3500 Average Reward: 5.94
Episode 4000 Average Reward: 6.01
Episode 4500 Average Reward: 6.04
Episode 5000 Average Reward: 6.0
5000
CPU times: user 1.68 s, sys: 117 ms, total: 1.8 s
Wall time: 1.69 s

```



Setting up environment
Running random simulation
Resetting environment



DEEP REINFORCEMENT LEARNING

CODE

```
#----- BRAIN -----

class Brain:

    """The 'brain' of the agent, where the model is created and held.

    state_dim (int): the size of the observation space
    action_dim (int): the size of the action space

    """

    def __init__(self, state_dim, action_dim, weights=None):
        self.state_dim = state_dim
        self.action_dim = action_dim

        self.model = self._createModel()

        if weights:
            self.model.load_weights("brain.h5")

    def _createModel(self):
        # Creates a Sequential Keras model
        # This acts as the Deep Q-Network (DQN)

        model = Sequential()

        ### START CODE HERE ### (~ 3 lines of code)

        # 'Dense' is the basic form of a neural network layer

        # Input Layer with activation function relu and Hidden Layer with 1
        28 nodes
```



```

model.add(Dense(128, input_dim=self.state_dim, activation='relu'))
#Second Hidden layer with 128 nodes
model.add(Dense(128, activation='relu'))
#Output layer with activation linear.
#action_size=4
model.add(Dense(self.action_dim, activation='linear'))

### END CODE HERE ###

opt = RMSprop(lr=0.00025)
model.compile(loss='mse', optimizer=opt)

return model

def train(self, x, y, epoch=1, verbose=0):
    self.model.fit(x, y, batch_size=64, epochs=epoch, verbose=verbose)

def predict(self, s):
    return self.model.predict(s)

def predictOne(self, s):
    return self.predict(s.reshape(1, self.state_dim)).flatten()

#----- MEMORY -----
class Memory:    # stored as ( s, a, r, s_ )
    """The agent's 'memory', where experiences are stored
    """

    def __init__(self, capacity):

```

```

self.capacity = capacity
self.samples = []

def add(self, sample):
    # a sample should be an array [s, a, r, s_]
    # s: current state
    # a: current action
    # r: current reward
    # s_: next state
    self.samples.append(sample)

    if len(self.samples) > self.capacity:
        self.samples.pop(0)

def sample(self, n):
    n = min(n, len(self.samples))
    return random.sample(self.samples, n)

#----- AGENT -----
import math

class Agent:
    """The agent, which learns to navigate the environment

    """

    def __init__(self, state_dim, action_dim, memory_capacity = 10000,
                  batch_size = 64, gamma = 0.99, lamb = 0.001,
                  max_epsilon = 1., min_epsilon = 0.01):
        self.state_dim = state_dim
        self.action_dim = action_dim

```

```

self.batch_size = batch_size

self.gamma = gamma # discount rate, to calculate the future discounted reward

self.lamb = lamb

self.max_epsilon = max_epsilon
self.epsilon = max_epsilon
self.min_epsilon = min_epsilon


self.brain = Brain(state_dim, action_dim)
self.memory = Memory(memory_capacity)

self.steps = 0
self.epsilons = []


def act(self, s, verbose=False):
    """The policy of the agent:

    Here, we determine if we explore (take a random action) based on epsilon.

    If not, we have the model predict the Q-Values for the state,
    then take the action which maximizes those values.

    """
    if random.random() < self.epsilon:
        if verbose:
            print("Random Action.")

        return random.randint(0, self.action_dim-1)
    else:
        actions = self.brain.predictOne(s)

        if verbose:
            print("Actions:", actions)

        return np.argmax(actions)

```

```

def observe(self, sample): # in (s, a, r, s_) format
    """The agent observes an event.

    We pass a sample (state, action, reward, next state) to be stored i
n memory.

    We then increment the step count and adjust epsilon accordingly.
    """
    self.memory.add(sample)

    # slowly decrease Epsilon based on our eperience
    self.steps += 1

    ### START CODE HERE ### (~ 1 line of code)

    self.epsilon=self.min_epsilon+(self.max_epsilon-
self.min_epsilon)* math.exp((-self.lamb)*abs(self.steps))

    # $\epsilon = \epsilon_{min} + (\epsilon_{max} - \epsilon_{min}) * e^{-\lambda |S|}$ 

    ### END CODE HERE ###

    self.epsilons.append(self.epsilon)

def replay(self):
    """The agent learns based on previous experiences.

    We sample observations (state, action, reward, next state) from mem
ory.

    We train the model based on these observations.
    """

    # Random sample of experiences
    batch = self.memory.sample(self.batch_size)
    batch_size = len(batch)

```

```

# Extracting states ('current' and 'next') from samples
no_state = np.zeros(self.state_dim)

states = np.array([ o[0] for o in batch ])

states_next = np.array([ (no_state if o[3] is None else o[3]) for o
in batch ])


# Estimating Q-Values for states
q_vals = self.brain.predict(states)
q_vals_next = self.brain.predict(states_next)


# Setting up training data
x = np.zeros((batch_size, self.state_dim))
y = np.zeros((batch_size, self.action_dim))
done=False

for i in range(batch_size):
    obs = batch[i]
    st = obs[0];
    act = obs[1];
    rew = obs[2];
    st_next = obs[3]
    t = q_vals[i]

    ### START CODE HERE ### (~ 4 line of code)

    if st_next is None:
        t[act]=rew
    else:
        t[act] = (rew + self.gamma *np.amax(q_vals_next[i]))

```

```

    ### END CODE HERE ###

    # Set training data

    x[i] = st

    y[i] = t

    # Train
    self.brain.train(x, y)

#----- MAIN -----
print('Setting up environment')
env = Environment(5)

state_dim = 4
action_dim = 4 # left, right, up, down
print('Setting up agent')
MAX_EPSILON = 1 # the rate in which an agent randomly decides its action
MIN_EPSILON = 0.05 # min rate in which an agent randomly decides its action
LAMBDA = 0.00005 # speed of decay for epsilon
num_episodes = 10000 # number of games we want the agent to play

VERBOSE = False
agent = Agent(state_dim, action_dim, lamb=LAMBDA,
              max_epsilon=MAX_EPSILON, min_epsilon=MIN_EPSILON)
env.reset()
episode_rewards = []
epsilons = []
t0 = time.time()
frames = []

```

```

print('Running simulation')

for episode in range(num_episodes):
    s = env.reset() # Initial state
    if episode % 1000 == 0:
        fig = env.render()
        frames.append(fig)
    R = 0
    while True:
        a = agent.act(s, verbose=VERBOSE)

        s_, r, done = env.step(a)

        if done: # terminal state
            s_ = None

        agent.observe( (s, a, r, s_) )
        agent.replay()

        s = s_
        R += r

    if episode % 1000 == 0:
        fig = env.render()
        frames.append(fig)

    if VERBOSE:
        print("Action:", a)
        print("Reward:", r)

    if done:

```

```
        break

    epsilons.append(agent.epsilon)
    episode_rewards.append(R)

    if episode % 100 == 0:
        print('Episode {}'.format(episode))
        print('Time Elapsed: {0:.2f}s'.format(time.time() - t0))
        print('Epsilon {}'.format(epsilons[-1]))
        print('Last Episode Reward: {}'.format(R))
        print('Episode Reward Rolling Mean: {}'.format(np.mean(episode_rewards[:-100])))
        print('-'*10)

agent.brain.model.save("brain.h5")
```

OUTPUT


```
Setting up environment
Setting up agent
Running simulation
/usr/local/lib/python3.7/dist-packages/keras/optimizer_v2/rmsprop.py:130: User
  super(RMSprop, self).__init__(name, **kwargs)
/usr/local/lib/python3.7/dist-packages/keras/engine/training_v1.py:2079: User
  updates=self.state_updates,
Episode 0
Time Elapsed: 1.91s
Epsilon 0.9995251187302109
Last Episode Reward: 0
Episode Reward Rolling Mean: nan
-----
/usr/local/lib/python3.7/dist-packages/numpy/core/fromnumeric.py:3373: Runtime
  out=out, **kwargs)
/usr/local/lib/python3.7/dist-packages/numpy/core/_methods.py:170: RuntimeWarn
  ret = ret.dtype.type(ret / rcount)
Episode 100
Time Elapsed: 12.45s
Epsilon 0.9532162322387105
Last Episode Reward: 6
Episode Reward Rolling Mean: 0.0
-----
Episode 200
Time Elapsed: 23.07s
Epsilon 0.9091658567921319
Last Episode Reward: 0
Episode Reward Rolling Mean: 0.2376237623762376
-----
Episode 300
Time Elapsed: 33.76s
Epsilon 0.8673455739778486
Last Episode Reward: 0
Episode Reward Rolling Mean: 0.527363184079602
-----
```

Episode 9400
Time Elapsed: 1025.64s
Epsilon 0.06291658293417288
Last Episode Reward: 7
Episode Reward Rolling Mean: 6.281367594882271



Episode 9500
Time Elapsed: 1036.22s
Epsilon 0.06237108611129126
Last Episode Reward: 7
Episode Reward Rolling Mean: 6.295394107009892

Episode 9600
Time Elapsed: 1046.45s
Epsilon 0.06185218198747368
Last Episode Reward: 8
Episode Reward Rolling Mean: 6.307546574044838

Episode 9700
Time Elapsed: 1056.87s
Epsilon 0.061358450275701804
Last Episode Reward: 8
Episode Reward Rolling Mean: 6.320487449224039

Episode 9800
Time Elapsed: 1067.66s
Epsilon 0.06087766913536543
Last Episode Reward: 8
Episode Reward Rolling Mean: 6.335429337181734

Episode 9900
Time Elapsed: 1078.60s
Epsilon 0.060408387636363046
Last Episode Reward: 8
Episode Reward Rolling Mean: 6.3457810427507395

COMPARISON

MOUNTAIN CAR

	<u>Time taken</u>	<u>Episode count</u>	<u>Average Reward</u>
Reinforcement Learning	111 seconds	5000	-161.02
Deep Reinforcement Learning	606.494 seconds	250	-200

ROULETTE

	<u>Time taken</u>	<u>Episode count</u>	<u>Average Reward</u>
Reinforcement Learning	43 seconds	50000	0
Deep Reinforcement Learning	568.864 seconds	558	-11.2

CAR RACING

	<u>Time taken</u>	<u>Step count</u>	<u>Average Reward</u>
<u>Deep Reinforcement Learning</u>	1159.031 seconds	10000	-81.618

USER INPUT GRAPH

	<u>Time taken</u>	<u>Step count</u>	<u>Average Reward</u>
Reinforcement Learning	1.69 seconds	5000	6
Deep Reinforcement Learning	1078.60 seconds	10000	6.345781027